

## ▼ Churn Modelling Project

This data set contains details of a bank's customers and the target variable is a binary variable reflecting the fact whether the customer left the bank

(closed his account) or he continues to be a customer.

```
# import basic libraries
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings('ignore')

# TO load dataset
df=pd.read_csv('/content/drive/MyDrive/Palak Deep Learning/Churn_Modelling.csv')
```

```
# To show first 5 record
df.head()
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure
0	1	15634602	Hargrave	619	France	Female	42	2
1	2	15647311	Hill	608	Spain	Female	41	1
2	3	15619304	Onio	502	France	Female	42	8
3	4	15701354	Boni	699	France	Female	39	1
4	5	15737888	Mitchell	850	Spain	Female	43	2

```
# To show the numbers of rows and columns
df.shape
```

```
(10000, 14)
```

```
# To delete unwanted features permanently in dataset
df.drop(["RowNumber","CustomerId","Surname"],axis=1,inplace=True)
```

```
# TO show the null values
df.isnull().sum()
```

```


CreditScore      0
Geography        0
Gender           0
Age              0
Tenure           0
Balance          0
NumOfProducts   0
HasCrCard        0
IsActiveMember   0
EstimatedSalary  0
Exited           0
dtype: int64

```

```

# To visualize the null value
sns.heatmap(df.isnull())
plt.show()

```



```

# To check the duplicates rows
df.duplicated().sum()

0

```

```

# To show the datatypes
df.dtypes

```

```

CreditScore      int64
Geography        object
Gender           object
Age              int64
Tenure           int64
Balance          float64
NumOfProducts    int64
HasCrCard        int64
IsActiveMember    int64
EstimatedSalary  float64

```

```
Exited          int64
dtype: object
```

```
# create 1st datafram df_cat to hold the categorical data
df_cat=df.select_dtypes(object)
```

```
# creating 2nd dataframe df_num to hold numeric data
df_num=df.select_dtypes(['int64','float64'])
```

```
# converting all the categorical data into numeric data
# we use labelEncoder to convert the data
from sklearn.preprocessing import LabelEncoder
# create object of Label Encoder
for col in df_cat:
    le=LabelEncoder()
    df_cat[col]=le.fit_transform(df_cat[[col]])
```

```
# concatenating of both the dataframe df_cat and df_num and hold into df_new
df_new=pd.concat([df_cat,df_num],axis=1)
```

```
# New dataframe
df_new.head()
```

	Geography	Gender	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember
0	0	0	619	42	2	0.00	1	1	
1	2	0	608	41	1	83807.86	1	0	
2	0	0	502	42	8	159660.80	3	1	
3	0	0	699	39	1	0.00	2	0	
4	2	0	850	43	2	125510.82	1	1	

```
df_new.dtypes
```

```
Geography          int64
Gender             int64
CreditScore        int64
Age                int64
Tenure             int64
Balance            float64
NumOfProducts      int64
HasCrCard          int64
IsActiveMember     int64
EstimatedSalary    float64
Exited             int64
dtype: object
```

```
# first check how many samples in both classes
```

```
df['Exited'].value_counts()
```

```
0    7963
```

```
1    2037
```

```
Name: Exited, dtype: int64
```

Here as we can see the data is not equally distributed. we have to divide the data into equally part

```
# Visualization the unbalance data
```


```
plt.figure(figsize=(10,5))
```

```
sns.countplot(data=df,x='Exited')
```

```
f = df['Exited'].value_counts()
```

```
plt.yticks(f)
```

```
plt.show()
```



```
# select input and output
```

```
X=df_new.drop("Exited",axis=1) # input
```

```
Y=df_new["Exited"] # output
```

```
# Train Test split the data
```

```
from sklearn.model_selection import train_test_split
```

```
X_train ,X_test,Y_train,Y_test=train_test_split(X,Y,test_size=0.3,random_state=1)
```

```
## Apply scaling (Standard Scaler)
```

```
from sklearn.preprocessing import StandardScaler
```

```
# create object of StandardScaler class
```

```
ss=StandardScaler()
```

```
X_train=ss.fit_transform(X_train)
```

```
X_test=ss.transform(X_test)
```

```
X.shape
```

```
(10000, 10)
```

## Applying 'RandomOverSampler' to make data equal.

From the above value\_counts of 'Exited' Column and from the Graph There are about: '7963' out of '2037'  
Which means the Data is not properly distributed we have to divided the data qually part then, use  
RandomOverSampler First import RandomOverSampler

```
# Apply random over sampling : inbuild class
from imblearn.over_sampling import RandomOverSampler
```

```
# Create the object of RandomOverSampler
ros = RandomOverSampler(random_state=1)
```

```
# Before apply RandomOverSampler on training data
Y_train.value_counts()
```

```
0    5590
1    1410
Name: Exited, dtype: int64
```

```
# Applied OverSampler on Training data (70%)
X_train_ros,Y_train_ros = ros.fit_resample(X_train,Y_train)
```

```
# Check after apply RandomOverSampler
Y_train_ros.value_counts()
```

```
0    5590
1    5590
Name: Exited, dtype: int64
```

```
# Before apply RandomOverSampler on testing data
Y_test.value_counts()
```

```
0    2373
1     627
Name: Exited, dtype: int64
```

```
# Also apply RandomOverSampler on tesing data (30%)
X_test_ros,Y_test_ros = ros.fit_resample(X_test,Y_test)
```

```
Y_test_ros.value_counts()
```

```
0    2373
1    2373
Name: Exited, dtype: int64
```

Apply Neural Network

```

#First step when we take Neurons only 10,. means Total Number of Neurons = Total Nu
# create a neural network
import tensorflow as tf
# create object of sequential class
model=tf.keras.Sequential([
    tf.keras.layers.Dense(units=10,activation='relu',input_shape=(X.shape[1],)),
    # hidden layer 1
    tf.keras.layers.Dense(units=10,activation='relu'), # second hidden layer
    tf.keras.layers.Dense(units=1,activation='sigmoid') # output layer
])

```

```
model.summary()
```

```
Model: "sequential"
```

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 10)	110
dense_1 (Dense)	(None, 10)	110
dense_2 (Dense)	(None, 1)	11
Total params: 231		
Trainable params: 231		
Non-trainable params: 0		

```
# Compile the model
```

```
model.compile(optimizer='Adam',loss='binary_crossentropy',metrics=['accuracy'])
```

```

#Create Early Stopping means create a call back
from tensorflow.keras.callbacks import EarlyStopping
callback=EarlyStopping(
    monitor="val_loss",
    min_delta=0.00001,
    patience=20,
    verbose=1,
    mode="auto",
    baseline=None,
    restore_best_weights=False
)

```

```
# Train the Model
```

```
trained_model = model.fit(X_train_ros,Y_train_ros,batch_size=20,epochs=3500,validation
```

```
Epoch 1/3500
```

```
559/559 [=====] - 2s 2ms/step - loss: 0.6072 - accuracy: 0.6727
```

```
Epoch 2/3500
```

```
559/559 [=====] - 1s 2ms/step - loss: 0.5385 - accuracy: 0.7296
```

```
Epoch 3/3500
```

```

559/559 [=====] - 2s 4ms/step - loss: 0.5086 - accuracy: 0.7497
Epoch 4/3500
559/559 [=====] - 2s 4ms/step - loss: 0.4914 - accuracy: 0.7589
Epoch 5/3500
559/559 [=====] - 2s 4ms/step - loss: 0.4803 - accuracy: 0.7654
Epoch 6/3500
559/559 [=====] - 2s 4ms/step - loss: 0.4736 - accuracy: 0.7694
Epoch 7/3500
559/559 [=====] - 2s 4ms/step - loss: 0.4683 - accuracy: 0.7742
Epoch 8/3500
559/559 [=====] - 3s 5ms/step - loss: 0.4643 - accuracy: 0.7758
Epoch 9/3500
559/559 [=====] - 3s 5ms/step - loss: 0.4609 - accuracy: 0.7786
Epoch 10/3500
559/559 [=====] - 2s 4ms/step - loss: 0.4585 - accuracy: 0.7817
Epoch 11/3500
559/559 [=====] - 3s 5ms/step - loss: 0.4558 - accuracy: 0.7856
Epoch 12/3500
559/559 [=====] - 2s 4ms/step - loss: 0.4536 - accuracy: 0.7841
Epoch 13/3500
559/559 [=====] - 2s 4ms/step - loss: 0.4513 - accuracy: 0.7876
Epoch 14/3500
559/559 [=====] - 2s 4ms/step - loss: 0.4504 - accuracy: 0.7849
Epoch 15/3500
559/559 [=====] - 2s 3ms/step - loss: 0.4488 - accuracy: 0.7861
Epoch 16/3500
559/559 [=====] - 1s 2ms/step - loss: 0.4470 - accuracy: 0.7889
Epoch 17/3500
559/559 [=====] - 1s 2ms/step - loss: 0.4463 - accuracy: 0.7887
Epoch 18/3500
559/559 [=====] - 1s 2ms/step - loss: 0.4445 - accuracy: 0.7907
Epoch 19/3500
559/559 [=====] - 1s 2ms/step - loss: 0.4443 - accuracy: 0.7911
Epoch 20/3500
559/559 [=====] - 1s 2ms/step - loss: 0.4432 - accuracy: 0.7913
Epoch 21/3500
559/559 [=====] - 1s 2ms/step - loss: 0.4424 - accuracy: 0.7927
Epoch 22/3500
559/559 [=====] - 1s 2ms/step - loss: 0.4420 - accuracy: 0.7897
Epoch 23/3500
559/559 [=====] - 1s 2ms/step - loss: 0.4411 - accuracy: 0.7934
Epoch 24/3500
559/559 [=====] - 1s 2ms/step - loss: 0.4407 - accuracy: 0.7903
Epoch 25/3500
559/559 [=====] - 1s 2ms/step - loss: 0.4402 - accuracy: 0.7928
Epoch 26/3500
559/559 [=====] - 1s 2ms/step - loss: 0.4391 - accuracy: 0.7929
Epoch 27/3500
559/559 [=====] - 1s 2ms/step - loss: 0.4388 - accuracy: 0.7936
Epoch 28/3500
559/559 [=====] - 1s 2ms/step - loss: 0.4388 - accuracy: 0.7936
Epoch 29/3500


```

```
# Visualisation Training and Testing Loss
```

```


plt.plot(trained_model.history['loss'],label='Training Loss')
plt.plot(trained_model.history['val_loss'],label='Testing Loss')
plt.legend()
plt.show()

```



# Visualise

```
plt.plot(trained_model.history['accuracy'],label='Training Accuracy')
plt.plot(trained_model.history['val_accuracy'],label='Testing Accuracy')
plt.legend()
plt.show()
```



#Here we can see gap between Accuracy and loss, Which shows the model is Overfit

## Conclusion

To reduce overfitting we will use Regularization: Regularization is a technique to prevent the model from overfitting. There are two type of technique of Regularization 1.Ridge L2 2.Lasso L1 here we use Ridge L2 Regularization .

```
# To create overfitting, use regularation : Ridge means L2
# create a neural network
# Create object of Sequential Class
# To reduce overfitting , use Regularation : Ridge means L2
# Create a Neural Network
```

```
from keras import regularizers
model = tf.keras.models.Sequential([
    tf.keras.layers.Dense(units=10,activation='relu',input_shape=(X.shape[1],),kernel_regularizer=regularizers.L2(0.01)),
    tf.keras.layers.Dense(units=10,activation='relu'),
    tf.keras.layers.Dense(units=1,activation='sigmoid',kernel_regularizer=regularizers.L2(0.01))
])
```



```
# Compile the model
```

```
model.compile(optimizer='Adam',loss='binary_crossentropy',metrics=['accuracy'])
```

```
#Train the Model
```

```
trained_model = model.fit(X_train_ros,Y_train_ros,batch_size=20,epochs=3500,validation_data=(X_val_ros,Y_val_ros))
```

```
Epoch 1/3500
```

```
559/559 [=====] - 2s 2ms/step - loss: 0.7193 - accuracy: 0.6258
```

```
Epoch 2/3500
```

```
559/559 [=====] - 1s 2ms/step - loss: 0.6057 - accuracy: 0.7202
```

```
Epoch 3/3500
```

```
559/559 [=====] - 1s 2ms/step - loss: 0.5590 - accuracy: 0.7436
```

```
Epoch 4/3500
```

```
559/559 [=====] - 1s 2ms/step - loss: 0.5370 - accuracy: 0.7514
```

```
Epoch 5/3500
```

```
559/559 [=====] - 1s 2ms/step - loss: 0.5275 - accuracy: 0.7541
```

```
Epoch 6/3500
```

```
559/559 [=====] - 1s 2ms/step - loss: 0.5218 - accuracy: 0.7577
```

```
Epoch 7/3500
```

```
559/559 [=====] - 1s 2ms/step - loss: 0.5171 - accuracy: 0.7588
```

```
Epoch 8/3500
```

```
559/559 [=====] - 1s 2ms/step - loss: 0.5146 - accuracy: 0.7618
```

```
Epoch 9/3500
```

```
559/559 [=====] - 1s 2ms/step - loss: 0.5123 - accuracy: 0.7595
```

```
Epoch 10/3500
```

```
559/559 [=====] - 1s 2ms/step - loss: 0.5103 - accuracy: 0.7601
```

```
Epoch 11/3500
```

```
559/559 [=====] - 1s 2ms/step - loss: 0.5082 - accuracy: 0.7625
```

```
Epoch 12/3500
```

```
559/559 [=====] - 1s 2ms/step - loss: 0.5066 - accuracy: 0.7645
```

```
Epoch 13/3500
```

```
559/559 [=====] - 1s 2ms/step - loss: 0.5057 - accuracy: 0.7636
```

```
Epoch 14/3500
```

```
559/559 [=====] - 1s 2ms/step - loss: 0.5038 - accuracy: 0.7623
```

```
Epoch 15/3500
```

```
559/559 [=====] - 1s 2ms/step - loss: 0.5025 - accuracy: 0.7628
```

```
Epoch 16/3500
```

```
559/559 [=====] - 1s 2ms/step - loss: 0.5022 - accuracy: 0.7656
```

```
Epoch 17/3500
```

```
559/559 [=====] - 1s 2ms/step - loss: 0.5012 - accuracy: 0.7653
```

```
Epoch 18/3500
```

```
559/559 [=====] - 1s 2ms/step - loss: 0.5004 - accuracy: 0.7659
```

```
Epoch 19/3500
```

```
559/559 [=====] - 1s 2ms/step - loss: 0.4992 - accuracy: 0.7643
```

```
Epoch 20/3500
```

```
559/559 [=====] - 1s 2ms/step - loss: 0.4984 - accuracy: 0.7648
```

```
Epoch 21/3500
```

```
559/559 [=====] - 1s 2ms/step - loss: 0.4972 - accuracy: 0.7676
```

```
Epoch 22/3500
```

```
559/559 [=====] - 2s 4ms/step - loss: 0.4966 - accuracy: 0.7686
```

```
Epoch 23/3500
```

```
559/559 [=====] - 2s 4ms/step - loss: 0.4962 - accuracy: 0.7655
```

```
Epoch 24/3500
```

```
559/559 [=====] - 1s 2ms/step - loss: 0.4956 - accuracy: 0.7682
```

```
Epoch 25/3500
```

```
559/559 [=====] - 1s 3ms/step - loss: 0.4948 - accuracy: 0.7652
```

```
Epoch 26/3500
```

```
559/559 [=====] - 2s 4ms/step - loss: 0.4951 - accuracy: 0.7668
```

```
Epoch 27/3500
```

```
559/559 [=====] - 2s 3ms/step - loss: 0.4939 - accuracy: 0.7676
```


```
Epoch 28/3500
559/559 [=====] - 1s 2ms/step - loss: 0.4935 - accuracy: 0.7684
Epoch 29/3500
```

```
print("Training Loss and Training Accuracy:",model.evaluate(X_train_ros,Y_train_ros))
print("Testing Loss and Testing Accuracy:",model.evaluate(X_train_ros,Y_train_ros))
```

```
350/350 [=====] - 0s 1ms/step - loss: 0.4691 - accuracy: 0.7841
Training Loss and Training Accuracy: [0.4690917432308197, 0.7840787172317505]
350/350 [=====] - 0s 1ms/step - loss: 0.4691 - accuracy: 0.7841
Testing Loss and Testing Accuracy: [0.4690917432308197, 0.7840787172317505]
```


```
# Visualisation Training and Testing Loss
```

```
plt.plot(trained_model.history['loss'],label='Training Loss')
plt.plot(trained_model.history['val_loss'],label='Testing Loss')
plt.legend()
plt.show()
```



```
# Visualise
```

```
plt.plot(trained_model.history['accuracy'],label='Training Accuracy')
plt.plot(trained_model.history['val_accuracy'],label='Testing Accuracy')
plt.legend()
plt.show()
```



## Conclusion

Now After using Regularization we got good accuracy, but we will still try to get more accuracy.

Then now we will increase the numbers of neurons and if we increase the number of neurons we will use dropout . Dropout is a technique that drops the number of neurons from the neural network or 'ignores' them during training.

Its means dropout handle the huge amount of number of neurons.

Dropout always 20,30,and 50% not more than 50%

# Create a Neural Network

```
from keras import regularizers
from keras.layers import Dropout
model = tf.keras.models.Sequential([
    tf.keras.layers.Dense(units=1000,activation='relu',input_shape=(X.shape[1],)),
    tf.keras.layers.Dense(units=1000,activation='relu',kernel_regularizer=regularizers.l2(0.01)),
    tf.keras.layers.Dense(units=1,activation='sigmoid',kernel_regularizer=regularizers.l2(0.01))
])
```

# Compile the model

```
model.compile(optimizer='adam',loss='binary_crossentropy',metrics=['accuracy'])
```

# Train the Model

```
trained_model= model.fit(X_train_ros,Y_train_ros,batch_size=32,epochs=3500,
                        validation_data=(X_test_ros,Y_test_ros),callbacks=callback)
```

```
Epoch 1/3500
350/350 [=====] - 7s 18ms/step - loss: 1.2342 - accuracy: 0.721
Epoch 2/3500
350/350 [=====] - 6s 17ms/step - loss: 0.6222 - accuracy: 0.742
Epoch 3/3500
350/350 [=====] - 6s 17ms/step - loss: 0.6106 - accuracy: 0.745
Epoch 4/3500
350/350 [=====] - 6s 16ms/step - loss: 0.6103 - accuracy: 0.738
Epoch 5/3500
350/350 [=====] - 6s 16ms/step - loss: 0.6069 - accuracy: 0.742
Epoch 6/3500
350/350 [=====] - 6s 16ms/step - loss: 0.6033 - accuracy: 0.742
Epoch 7/3500
350/350 [=====] - 6s 17ms/step - loss: 0.6037 - accuracy: 0.743
Epoch 8/3500
350/350 [=====] - 6s 16ms/step - loss: 0.6043 - accuracy: 0.742
Epoch 9/3500
350/350 [=====] - 6s 16ms/step - loss: 0.6014 - accuracy: 0.743
Epoch 10/3500
350/350 [=====] - 6s 17ms/step - loss: 0.6003 - accuracy: 0.742
Epoch 11/3500
350/350 [=====] - 6s 17ms/step - loss: 0.5999 - accuracy: 0.743
Epoch 12/3500
350/350 [=====] - 6s 17ms/step - loss: 0.6018 - accuracy: 0.744
Epoch 13/3500
350/350 [=====] - 6s 16ms/step - loss: 0.5993 - accuracy: 0.746
Epoch 14/3500
350/350 [=====] - 6s 16ms/step - loss: 0.5993 - accuracy: 0.743
```

```

Epoch 15/3500
350/350 [=====] - 6s 17ms/step - loss: 0.6006 - accuracy: 0.742
Epoch 16/3500
350/350 [=====] - 6s 16ms/step - loss: 0.6002 - accuracy: 0.743
Epoch 17/3500
350/350 [=====] - 6s 16ms/step - loss: 0.5991 - accuracy: 0.748
Epoch 18/3500
350/350 [=====] - 6s 16ms/step - loss: 0.5990 - accuracy: 0.747
Epoch 19/3500
350/350 [=====] - 6s 16ms/step - loss: 0.5965 - accuracy: 0.744
Epoch 20/3500
350/350 [=====] - 6s 17ms/step - loss: 0.5975 - accuracy: 0.745
Epoch 21/3500
350/350 [=====] - 7s 19ms/step - loss: 0.5989 - accuracy: 0.743
Epoch 22/3500
350/350 [=====] - 6s 18ms/step - loss: 0.5992 - accuracy: 0.742
Epoch 23/3500
350/350 [=====] - 6s 16ms/step - loss: 0.5959 - accuracy: 0.746
Epoch 24/3500
350/350 [=====] - 6s 16ms/step - loss: 0.5968 - accuracy: 0.745
Epoch 25/3500
350/350 [=====] - 6s 17ms/step - loss: 0.5964 - accuracy: 0.746
Epoch 26/3500
350/350 [=====] - 6s 17ms/step - loss: 0.5964 - accuracy: 0.747
Epoch 27/3500
350/350 [=====] - 6s 17ms/step - loss: 0.5985 - accuracy: 0.744
Epoch 28/3500
350/350 [=====] - 6s 17ms/step - loss: 0.5972 - accuracy: 0.744
Epoch 29/3500

```

```

print("Training Loss and Training Accuracy:",model.evaluate(X_train_ros,Y_train_ros))
print("Testing Loss and Testing Accuracy:",model.evaluate(X_train_ros,Y_train_ros))

```

```

350/350 [=====] - 2s 7ms/step - loss: 0.5902 - accuracy: 0.7485
Training Loss and Training Accuracy: [0.5901646018028259, 0.748479425907135]
350/350 [=====] - 1s 4ms/step - loss: 0.5902 - accuracy: 0.7485
Testing Loss and Testing Accuracy: [0.5901646018028259, 0.748479425907135]


```

# Visualisation Training and Testing Loss

```

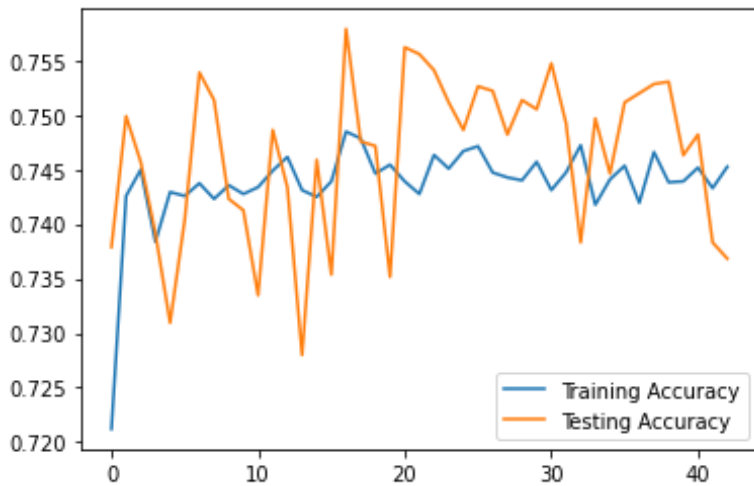
plt.plot(trained_model.history['loss'],label='Training Loss')
plt.plot(trained_model.history['val_loss'],label='Testing Loss')
plt.legend()
plt.show()

```



## # Visualisation Training and Testing Accuracy

```
plt.plot(trained_model.history['accuracy'],label='Training Accuracy')  
plt.plot(trained_model.history['val_accuracy'],label='Testing Accuracy')  
plt.legend()  
plt.show()
```



## Conclusion

We can see that,After applying Dropout we got better percentage as what we had applied using Regularization.